

UNIT 3: Storage Service

AWS S3 (Simple Storage Service) – Quick Overview 🚀

AWS S3 (Simple Storage Service) is a scalable, durable, and secure object storage service used for storing and retrieving data. It is designed for high availability (99.99%) and 11 nines (99.999999999%) durability.

Key Features of S3

- Object Storage – Stores data as objects (files + metadata) in buckets.
- Unlimited Storage – No total storage limit; individual objects can be up to 5 TB.
- Security & Access Control – Uses IAM policies, Bucket Policies, ACLs, and Encryption for protection.
- Versioning – Keeps multiple versions of an object for backup & recovery.
- Lifecycle Policies – Automatically moves data between storage classes or deletes objects.
- Replication – Cross-Region Replication (CRR) & Same-Region Replication (SRR) for redundancy.
- Static Website Hosting – Can serve websites directly from S3.

Amazon S3 Use cases

- Backup and storage
- Disaster Recovery
- Archive
- Hybrid Cloud storage
- Application hosting
- Media hosting
- Data lakes & big data analytics
- Static website

Amazon S3 - Objects

- Objects (files) have a Key
- The **key** is the FULL path:
 - s3://my-bucket/my_file.txt
 - s3://my-bucket/my_folder1/another_folder/my_file.txt
- The key is composed of **prefix** + **object name**
 - s3://my-bucket/my_folder1/another_folder/my_file.txt
- There's no concept of "directories" within buckets (although the UI will trick you to think otherwise)
- Just keys with very long names that contain slashes ("/")
- Object values are the content of the body:
 - Max. Object Size is 5TB (5000GB)
 - If uploading more than 5GB, must use "multi-part upload"
- Metadata (list of text key / value pairs – system or user metadata)
- Tags (Unicode key / value pair – up to 10) – useful for security / lifecycle
- Version ID (if versioning is enabled)

Amazon S3 – Security

- User-Based
 - IAM Policies – which API calls should be allowed for a specific user from IAM
- Resource-Based
 - Bucket Policies – bucket wide rules from the S3 console - allows cross account
 - Object Access Control List (ACL) – finer grain (can be disabled)
 - Bucket Access Control List (ACL) – less common (can be disabled)
- Note: an IAM principal can access an S3 object if
 - The user IAM permissions ALLOW it OR the resource policy ALLOWS it
 - AND there's no explicit DENY

- Encryption: encrypt objects in Amazon S3 using encryption keys

S3 Bucket Policies

- JSON based policies
 - Resources: buckets and objects
 - Effect: Allow / Deny
 - Actions: Set of API to Allow or Deny
 - Principal: The account or user to apply the policy to
- Use S3 bucket for policy to:
 - Grant public access to the bucket
 - Force objects to be encrypted at upload
 - Grant access to another account (Cross Account)


```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicRead",
      "Effect": "Allow",
      "Principal": "*",
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::examplebucket/*"
      ]
    }
  ]
}
```

Another Bucket Level Policy: which deny List/put/get objects to all the users except UserA

```
{
  "Version": "2012-10-17",
  "Id": "Policy1741273722082",
  "Statement": [
    {
      "Sid": "AllowUserA",
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::854844598681:user/userA"
      },
      "Action": [
        "s3:GetObject",
        "s3:ListBucket",
        "s3:PutObject"
      ],
      "Resource": [
        "arn:aws:s3:::shrikant-demo-123",
        "arn:aws:s3:::shrikant-demo-123/*"
      ]
    },
    {
      "Sid": "DenyAllOthers",
      "Effect": "Deny",
      "Principal": "*",
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::shrikant-demo-123",
        "arn:aws:s3:::shrikant-demo-123/*"
      ],
      "Condition": {
        "StringNotEquals": {
          "aws:PrincipalArn": "arn:aws:iam::854844598681:user/userA"
        }
      }
    }
  ]
}
```

Public Access (Not IAM User):

Example: Public Access - Use Bucket Policy



Access to IAM User

Example: User Access to S3 – IAM permissions




EC2 Accessing S3

Example: EC2 instance access - Use IAM Roles



Bucket settings for Block Public Access



- These settings were created to prevent company data leaks
- If you know your bucket should never be public, leave these on
- Can be set at the account level

Amazon S3 - Versioning

- You can version your files in Amazon S3
- It is enabled at the bucket level
- Same key overwrite will change the “version”: 1, 2, 3....
- It is best practice to version your buckets
 - Protect against unintended deletes (ability to restore a version)
 - Easy roll back to previous version
- Notes:
 - Any file that is not versioned prior to enabling versioning will have version “null”
 - Suspending versioning does not delete the previous versions

S3 Storage Classes

- S3 Standard – High performance, frequently accessed data.
- S3 Intelligent-Tiering – Auto-moves data between storage classes based on access patterns.
- S3 Standard-IA (Infrequent Access) – Cheaper for less accessed data.
- S3 One Zone-IA – Lower cost, but stored in one availability zone.
- S3 Glacier – For long-term archival storage (retrieval in minutes to hours).
- S3 Glacier Deep Archive – Lowest cost, retrieval in hours.

Pricing Model

- Pay-as-you-go model based on:
 - Storage used (GBs/TBs)
 - Data retrieval & transfer
 - Number of requests (PUT, GET, DELETE, etc.)

- Versioning
- Cross-region replication
- Life Cycle Management
- Security & Encryption

Static Web-hosting with S3 bucket: [Steps to Host a Static Website on S3](#)

Events configuration on S3 buckets - [Events configuration on S3 buckets](#)

Enabling Cross-Account Access to S3 Objects

- Enabling cross-account access for S3
- S3 Data management and backup using 3rd Party applications.
- S3 Cross-Account Access and Pre-Signed URLs

Storage Gateway:

AWS Storage Gateway is a **hybrid cloud storage service** that connects your **on-premises** applications with AWS **cloud storage** (S3, EBS, Glacier). It enables seamless **data transfer** between your local environment and AWS for backup, archiving, disaster recovery, and hybrid cloud workloads.

Key Features

- ✓ **Extend On-Prem Storage to AWS** (For backup, disaster recovery, archiving)
- ✓ **Low-Latency Access to Cloud Storage** (Local caching speeds up access)
- ✓ **AWS Integration** (Works with S3, EBS, Glacier, FSx, etc.)
- ✓ **Encryption & Security** (Data is encrypted in transit and at rest)

Exercise:

Restrict Access to an S3 Bucket Using IAM Policy

Scenario

An employee should have **read-only access** to an S3 bucket, but should **not be able to delete files**.

Steps to Solve

1. Create an IAM Policy

- Go to **IAM Console** → **Policies** → **Create Policy**
- Use the following policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "VisualEditor0",
      "Effect": "Allow",
      "Action": [
        "s3:GetObject",
        "s3:ListBucket"
      ],
      "Resource": [
        "arn:aws:s3:::shrikant-demo-123/*",
        "arn:aws:s3:::shrikant-demo-123"
      ]
    },
    {
      "Sid": "VisualEditor1",
      "Effect": "Deny",
      "Action": "s3:DeleteObject",
      "Resource": [
        "arn:aws:s3:::shrikant-demo-123/*",
        "arn:aws:s3:::shrikant-demo-123"
      ]
    }
  ]
}
```

Attach the Policy to an IAM User

- Go to **IAM Users** → Select user → **Attach Policy**

Note : You Cannot Attach an IAM Policy to the Root User

AWS does not allow attaching IAM policies directly to the **root user**. However, you can **prevent the root user from deleting objects** using **Service Control Policies (SCPs)** if your AWS account is part of **AWS Organizations**.

Enable Cross-Region Replication (CRR) for Disaster Recovery

Scenario

You need to replicate objects from an S3 bucket in **us-east-1** to another bucket in **us-west-1**.

Steps to Solve

1. Enable Versioning on Both Buckets

- Source bucket: my-source-bucket
- Destination bucket: my-destination-bucket

2. Set Up Replication Rule

- Go to **Management** → **Create Replication Rule**
- Choose **Destination Bucket**
- Select **IAM Role** or create a new one
- Enable **Replicate Existing Objects**

Use S3 Lifecycle Policies to Move Files to Glacier

Scenario

A company wants to archive objects older than **30 days** to **S3 Glacier**. Configure an S3 Lifecycle Rule.

Steps to Solve

1. Go to Lifecycle Configuration

- In **S3 Console**, select the bucket
- Navigate to **Management** → **Create Lifecycle Rule**

2. Define Rule

- Name: MoveToGlacierRule
- Select **All objects in the bucket**
- Set **Transition** → Move objects to **S3 Glacier after 30 days**
- Save the rule and test it after 30 days

Enable Cross-Region Replication (CRR) for Disaster Recovery

Scenario

You need to replicate objects from an S3 bucket in **us-east-1** to another bucket in

You need to replicate objects from an S3 bucket in **us-east-1** to another bucket in **us-west-1**.

Steps to Solve

1. **Enable Versioning on Both Buckets**
 - Source bucket: my-source-bucket
 - Destination bucket: my-destination-bucket
2. **Set Up Replication Rule**
 - Go to **Management** → **Create Replication Rule**
 - Choose **Destination Bucket**
 - Select **IAM Role** or create a new one
 - Enable **Replicate Existing Objects**

Configure S3 Event Notifications to Trigger a Lambda Function

Scenario

You need to **trigger a Lambda function** whenever a new file is uploaded to an S3 bucket.

Steps to Solve

1. **Create an S3 Event Notification**
 - Go to **S3 Console** → Select the bucket
 - Navigate to **Properties** → **Event Notifications**
 - Create a new event for PUT operations
 - Choose **Lambda Function** as the destination
2. **Create a Lambda Function**
 - Go to **Lambda Console** → Create a function
 - Use the following Python code to print object details:

```
import json
def lambda_handler(event, context):
    print("New object uploaded:", json.dumps(event, indent=4))
```

lambda_handler is the entry point function that AWS Lambda automatically invokes when the function is triggered.

- The function receives **two parameters**:
 1. **event** – Contains **data** about the event that triggered the Lambda function (e.g., an S3 file upload event, an API Gateway request, etc.).
 2. **context** – Provides **runtime metadata** (e.g., function name, memory limit, request ID).

AWS calls this function **whenever the Lambda is triggered** (e.g., when an object is uploaded to an S3 bucket).

3. `print("New object uploaded:", json.dumps(event, indent=4))`
 - **print()** logs messages to **Amazon CloudWatch Logs**, helping with debugging.
 - "New object uploaded:" is a **simple log message**.
 - `json.dumps(event, indent=4)`:
 - Converts the **event dictionary** to a **formatted JSON string**.
 - `indent=4` makes it **more readable** (pretty-printed format).
 - This allows you to **see details** about the event in CloudWatch Logs.

3. Test by Uploading a File

Whats get printed :

```
{
  "Records": [
    {
      "eventVersion": "2.1",
      "eventSource": "aws:s3",
      "eventName": "ObjectCreated:Put",
      "s3": {
        "bucket": {
          "name": "my-upload-bucket"
        },
        "object": {
          "key": "image.jpg",
          "size": 51234
        }
      }
    }
  ]
}
```

Output in CloudWatch Logs:

```
New object uploaded: {
  "Records": [
    {
      "eventVersion": "2.1",
      "eventSource": "aws:s3",
      "eventName": "ObjectCreated:Put",
      "s3": {
        "bucket": {
          "name": "my-upload-bucket"
        },
        "object": {
          "key": "image.jpg",
          "size": 51234
        }
      }
    }
  ]
}
```

Modified Code to Log the Uploader's Identity

modified code to log the uploader's identity
import json

```
def lambda_handler(event, context):
    # Extract S3 bucket and object key
    bucket_name = event["Records"][0]["s3"]["bucket"]["name"]
    object_key = event["Records"][0]["s3"]["object"]["key"]

    # Extract uploader identity (IAM user, role, or assumed identity)
    user_identity = event["Records"][0].get("userIdentity", {}).get("principalId",
"Unknown User")

    # Log message with uploader details
    print(f"New object uploaded to S3 bucket '{bucket_name}': '{object_key}' by {user_identity}")

    # Pretty-print full event for debugging
    print("Event Details:", json.dumps(event, indent=4))
```

Configure S3 to Require MFA Delete

Scenario

To prevent accidental file deletions, configure **MFA Delete** for an S3 bucket.

Steps to Solve

1. Enable MFA Delete (via AWS CLI)

- Run the following command:

```
aws s3api put-bucket-versioning --bucket my-secure-bucket --versioning-configuration Status=Enabled,MFADelete=Enabled --mfa "SERIAL_NUMBER MFA_CODE"
```

- Replace SERIAL_NUMBER and MFA_CODE with actual values

2. Test by Trying to Delete a File

Enable MFA for deleting objects(this can not be done via Console)

```
aws s3api put-bucket-versioning --bucket shrikant-123-123-123 --versioning-configuration Status=Enabled,MFADelete=Enabled --mfa "arn:aws:iam::854844598681:mfa/root-account-mfa-device 121221"
```

Replace serial number with MFA device ARN Number and CURRENT MFA CODE

Your bucket Name -> shrikant-123-123-123

```
aws s3api delete-object \
--bucket shrikant-123-123-123 \
--key appl_stock.csv \
--version-id 1NkKxoXSEnpy2xub3.0yFXGjEgXn0KqK \
--mfa "arn:aws:iam::854844598681:mfa/root-account-mfa-device 704869"
```